



Numpy

(num-pie, not num-pee)



```
import numpy as np
```

Numpy arrays

- Numpy is a Python language extension that provides a new object -- the **array** -- for heterogeneous datasets
- Arrays are collections of objects with the same type, different from Python lists
- Operations on arrays are fast and vectorized

Numpy arrays

- To create an array from scratch, pass a list in to the `np.array` function

Numpy arrays

- Arrays can be multidimensional
- Some useful attributes to introspect:
 - **ndim**: number of dimensions
 - **shape**: number of elements along each dim
 - **size**: total number of elements

Numpy arrays

- Array operations act element-wise
- **Try these with lists and arrays to compare:**
 - multiply a list/array by 5
 - add two lists/arrays
 - subtract two lists/arrays

Numpy arrays

- Array operations are fast
- Creating a list of squared numbers from 0 to 10000 in pure-Python and with numpy:
 - `[x**2 for x in range(10000)]`
 - `np.arange(10000)**2`
- Numpy is 70x faster

Array operations

- Can be sorted in place with the `.sort()` method
- `.max()`, `.min()`, `.mean()` are handy

Array operations

- Can reshape arrays and change dimensionality, as long as number of elements is conserved
- e.g., can reshape a 1D, 100 element list to a 2D, 10 x 10 array

Numpy functions


- All math functions in numpy act on arrays, and return arrays
- Other useful functions are:
 - `arange(start, stop, increment)`
 - `linspace(start, stop, number of elements)`
 - `zeros(shape)`, `ones(shape)`, `eye(shape)`
 - `sum(array)`, `mean(array)`, `median(array)`

Array indexing / slicing

- Arrays support the same slicing and indexing patterns as Python lists
- Since arrays can be multi-dimensional, can slice or index along each dimension with commas:
 - `my_array[:10, 1:6, 5]`


Advanced indexing

- You can also use integer numpy arrays as an index selection
- For example, let's say we have the following array, and we want to get the 0th, 3rd, and 4th items:



Boolean arrays

- You could do the same sub-selection of items using a **boolean array**
- A boolean array is a numpy array of **True**'s and **False**'s:



Boolean arrays

- Boolean arrays are most useful when generated using logical operators
- For example:

`arr` =

| | | | | | |
|---|---|----|----|----|----|
| 4 | 8 | 15 | 16 | 23 | 42 |
|---|---|----|----|----|----|

`arr > 8` →


| | | | | | |
|---|---|---|---|---|---|
| F | F | T | T | T | T |
|---|---|---|---|---|---|

Boolean arrays

- We can chain together selection expressions like $(arr > 8)$ using the bitwise operators
 - and $\&$
 - or $|$

$arr =$

| | | | | | |
|---|---|----|----|----|----|
| 4 | 8 | 15 | 16 | 23 | 42 |
|---|---|----|----|----|----|

$(arr > 8) \& (arr < 40)$ 

| | | | | | |
|---|---|---|---|---|---|
| F | F | T | T | T | F |
|---|---|---|---|---|---|

Boolean arrays

- Use case: let's say I want to select out all stars within a range of RA, plot them in one color, and plot all others in another color

```
idx = (ra > 11.1324) & (ra < 31.5134)
selected = ra[idx]
not_selected = ra[np.logical_not(idx)]
```

Structured arrays

- Numpy arrays have to have homogeneous data types
- What if you have a table of numbers where each column has a type, but each is different?

| Name | ID | Height | Active |
|-------------|-----------|---------------|---------------|
| “mulder” | 11605 | 6.0 | False |
| “scully” | 42115 | 5.5 | True |

Structured arrays

- The solution is to use a Numpy structured array

```
data = [("mulder", 11605, 6.0, False),  
        ("scully", 42115, 5.5, True)]  
arr = np.array(data, dtype=[("Name", str),  
                             ("ID", int),  
                             ("Height", float),  
                             ("Active", bool)])
```

| Name | ID | Height | Active |
|----------|-------|--------|--------|
| "mulder" | 11605 | 6.0 | False |
| "scully" | 42115 | 5.5 | True |

Structured arrays

- Indexing with the column names gets you the data for a particular column
- Using an integer gets data from a row

`arr["ID"] -> [11605, 42115]`

`arr[0] -> ("mulder", 11605, 6.0, True)`



Break?

Dimensionality reduction

- Sometimes you'll want to perform operations along a single axis or dimension in a numpy array
- Most functions that reduce dimensionality accept an `axis` keyword that will perform the reduction over the specified axis
- `sum(arr, axis=0)`

Broadcasting

- How numpy treats arrays with different shapes during arithmetic operations

Broadcasting

- We're happy with the idea of multiplying a scalar by an array:
- $5 * \text{array}([1,2,3,4,5]) = \text{array}(5,10,15,20,25)$
- That's like multiplying a 0D array by a 1D array
- You can think of it as taking the 0D array, copying it 5 times to produce an array with the same shape as the 1D array, then performing the operation element-wise

Broadcasting

- That was a 0D array times a 1D array
- What if we add one dimension to each?
 - The rule still holds

| | | | | | | |
|---|----|---|---|---|---|---|
| 1 | 10 | 2 | | 1 | 2 | 3 |
| | | | X | 4 | 5 | 6 |
| | | | | 7 | 8 | 9 |

Broadcasting

| | | | | | | |
|---|----|---|----------|---|---|---|
| 1 | 10 | 2 | X | 1 | 2 | 3 |
| 1 | 10 | 2 | | 4 | 5 | 6 |
| 1 | 10 | 2 | | 7 | 8 | 9 |

Broadcasting

| | | |
|---|----|----|
| 1 | 20 | 6 |
| 4 | 50 | 12 |
| 7 | 60 | 18 |